

Importing Libraries

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Loading Dataset

```
df=pd.read_csv(r'C:\Users\LENOVO\Downloads\Churn_Modelling.csv')
```

```
df.shape
```

```
(10000, 14)
```

```
df.head()
```

	Row Number	Customer Id	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
df.tail()
```

	Row Number	Customer Id	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	96270.64	0

	Row Number	Customer Id	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	101699.77	0
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	42085.58	1
9998	9999	15682355	Sabattini	772	Germany	Male	42	3	75075.31	2	1	0	92888.52	1
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	38190.78	0

Visualizations


1. Univariate

```
sns.distplot(df['EstimatedSalary'],hist=True)
```

C:\Users\LENOVO\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
<AxesSubplot:xlabel='EstimatedSalary', ylabel='Density'>
```




2. Bivariate

```
sns.relplot(x='Age',y='Balance',data=df)  
<seaborn.axisgrid.FacetGrid at 0x1fd190d2070>
```




```
sns.catplot(x='Geography',y='CreditScore',data=df)  
# for categorical data  
<seaborn.axisgrid.FacetGrid at 0x1fd194e8d90>
```




3. Multivariate

```
sns.relplot(x='Age',y='Balance',hue='Geography',data=df)  
<seaborn.axisgrid.FacetGrid at 0x1fd1956e580>
```



```
sns.boxplot(x='Exited',y='CreditScore',hue='Gender',data=df)  
# for categorical data  
<AxesSubplot:xlabel='Exited', ylabel='CreditScore'>
```



Descriptive Statistics

df.describe()

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCreditCard	IsActiveMember	EstimatedSalary	Exited
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402769
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.000000	0.000000	11.580000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000	0.000000	51002.100000	0.000000

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
50%	5000.50000	1.569074e+07	652.00000	37.00000	5.00000	97198.540000	1.000000	1.00000	1.000000	100193.915000	0.00000
75%	7500.25000	1.575323e+07	718.00000	44.00000	7.00000	127644.240000	2.000000	1.00000	1.000000	149388.247500	0.00000
max	10000.00000	1.581569e+07	850.00000	92.00000	10.00000	250898.090000	4.000000	1.00000	1.000000	199992.480000	1.00000

Handling the missing(null) values

```
df.isnull().any()
```

```
RowNumber      False
CustomerId      False
Surname         False
CreditScore     False
Geography       False
Gender          False
Age             False
Tenure          False
Balance         False
NumOfProducts  False
HasCrCard       False
IsActiveMember  False
EstimatedSalary False
Exited          False
```

```
dtype: bool
```

```
df.isnull().sum()
```

```
RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
```

```
dtype: int64
```

```
# Since no null values are found no need to handle
```


Split the data into dependent and independent variables

```
x=df.iloc[:,3:13].values
print(x.shape)
y=df.iloc[:,13:14].values
print(y.shape)

(10000, 10)
(10000, 1)
```


Finding and Replacing Outliers

```
df.skew()
```

C:\Users\LENOVO\AppData\Local\Temp\ipykernel_3336\1665899112.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
df.skew()
RowNumber      0.000000
CustomerId      0.001149
CreditScore    -0.071607
Age             1.011320
Tenure          0.010991
Balance        -0.141109
NumOfProducts  0.745568
HasCrCard      -0.901812
IsActiveMember -0.060437
EstimatedSalary 0.002085
Exited         1.471611
dtype: float64
sns.boxplot(df["Age"])
```

```
<AxesSubplot:xlabel='Age'>
```



```
q1 = df["Age"].describe()["25%"]
```

```
q3 = df["Age"].describe()["75%"]
```

```
iqr = q3-q1
```

```
l_b = q1 -(1.5*iqr)
```

```
u_b = q3 + (1.5*iqr)
```

```
df[df["Age"]<l_b]
```

Row Number	Customer Id	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
------------	-------------	---------	-------------	-----------	--------	-----	--------	---------	---------------	-----------	----------------	-----------------	--------

```
df[df["Age"]>u_b]
```

Row Number	Customer Id	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
58	59	15623944	Tien	Spain	Female	66	4	0.00	1	1	0	1643.11	1
85	86	15805254	Ndukaku	Spain	Female	75	10	0.00	2	1	1	114675.75	0
104	105	15804919	Dunabin	Spain	Female	65	1	0.00	1	1	1	177655.68	1

	Row Num ber	Cust omer Id	Surna me	Cred itSco re	Geo grap hy	Ge nd er	A g e	Te nu re	Bal anc e	NumO fProdu cts	Has CrC ard	IsActiv eMem ber	Estima tedSala ry	Ex ite d
1 5 8	159	1558 9975	Macle an	646	Fran ce	Fe ma le	7 3	6	972 59.2 5	1	0	1	104719 .66	0
1 8 1	182	1578 9669	Hsia	510	Fran ce	Ma le	6 5	2	0.00	2	1	1	48071. 61	0
...
9 7 5 3	9754	1570 5174	Chied ozie	656	Ger man y	Ma le	6 8	7	153 545. 11	1	1	1	186574 .68	0
9 7 6 5	9766	1577 7067	Thom as	445	Fran ce	Ma le	6 4	2	136 770. 67	1	0	1	43678. 06	0
9 8 3 2	9833	1581 4690	Chuk wujek wu	595	Ger man y	Fe ma le	6 4	2	105 736. 32	1	1	1	89935. 73	1
9 8 9 4	9895	1570 4795	Vagin	521	Fran ce	Fe ma le	7 7	6	0.00	2	1	1	49054. 10	0
9 9 3 6	9937	1565 3037	Parks	609	Fran ce	Ma le	7 7	1	0.00	1	0	1	18708. 76	0

359 rows × 14 columns

#Replace the outlier

```
outlier_list = list(df[df["Age"] > u_b]["Age"])
```

```
outlier_list
```

```
[66,  
75,  
65,  
73,  
65,  
72,  
67,
```

67,
79,
80,
68,
75,
66,
66,
70,
63,
72,
64,
64,
70,
67,
82,
63,
69,
65,
69,
64,
65,
74,
67,
66,
67,
63,
70,
71,
72,
67,
74,
76,
66,
63,
66,
68,
67,
63,
71,
66,
69,
73,
65,
66,
64,
69,
64,
77,
74,
65,
70,
67,
69,
67,
74,
69,
74,

74,
64,
63,
63,
70,
74,
65,
72,
77,
66,
65,
74,
88,
63,
71,
63,
64,
67,
70,
68,
72,
71,
66,
75,
67,
73,
69,
76,
63,
85,
67,
74,
76,
66,
69,
66,
72,
63,
71,
63,
74,
67,
72,
72,
66,
84,
71,
66,
63,
74,
69,
84,
67,
64,
68,
66,
77,

70,
67,
79,
67,
76,
73,
66,
67,
64,
73,
76,
72,
64,
71,
63,
70,
65,
66,
65,
80,
66,
63,
63,
63,
63,
66,
74,
69,
63,
64,
76,
75,
68,
69,
77,
64,
66,
74,
71,
67,
68,
64,
68,
70,
64,
75,
66,
64,
78,
65,
74,
64,
64,
71,
77,
79,
70,


81,
64,
68,
68,
63,
79,
66,
64,
70,
69,
71,
72,
66,
68,
63,
71,
72,
72,
64,
78,
75,
65,
65,
67,
63,
68,
71,
73,
64,
66,
71,
69,
71,
66,
76,
69,
73,
64,
64,
75,
73,
71,
72,
63,
67,
68,
73,
67,
64,
63,
92,
65,
75,
67,
71,
64,
66,

64,
66,
67,
77,
92,
67,
63,
66,
66,
68,
65,
72,
71,
76,
63,
67,
67,
66,
67,
63,
65,
70,
72,
77,
74,
72,
73,
77,
67,
71,
64,
72,
81,
76,
69,
68,
74,
64,
64,
71,
68,
63,
67,
63,
64,
76,
63,
63,
68,
67,
72,
70,
81,
67,
73,
66,
68,

71,
66,
63,
75,
69,
64,
69,
70,
71,
71,
66,
70,
63,
64,
65,
63,
67,
71,
67,
65,
66,
63,
73,
66,
64,
72,
71,
69,
67,
64,
81,
73,
63,
67,
74,
83,
69,
71,
78,
63,
70,
69,
72,
70,
63,
74,
80,
69,
72,
67,
76,
71,
67,
71,
78,
63,
63,

```
68,  
64,  
70,  
78,  
69,  
68,  
64,  
64,  
77,  
77]  
outlier_dict = {}.fromkeys(outlier_list,u_b)  
  
outlier_dict  
{66: 62.0,  
75: 62.0,  
65: 62.0,  
73: 62.0,  
72: 62.0,  
67: 62.0,  
79: 62.0,  
80: 62.0,  
68: 62.0,  
70: 62.0,  
63: 62.0,  
64: 62.0,  
82: 62.0,  
69: 62.0,  
74: 62.0,  
71: 62.0,  
76: 62.0,  
77: 62.0,  
88: 62.0,  
85: 62.0,  
84: 62.0,  
78: 62.0,  
81: 62.0,  
92: 62.0,  
83: 62.0}  
df["Age"] = df["Age"].replace(outlier_dict)  
sns.boxplot(df["Age"])
```

<AxesSubplot:xlabel='Age'>



```
df[df["Age"]>u_b]
```

Row Number	Customer Id	Sur name	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCreditCard	IsActiveMember	EstimatedSalary	Exited
------------	-------------	----------	-------------	-----------	--------	-----	--------	---------	---------------	---------------	----------------	-----------------	--------

Check for Categorical columns and perform encoding.

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
ct=ColumnTransformer([('oh',OneHotEncoder(),[1,2])],remainder='passthrough')
x=ct.fit_transform(x)
print(x.shape)

(10000, 13)
# saving the data
import joblib
joblib.dump(ct,"churnct.pkl")

['churnct.pkl']
```

Split the data into training and testing

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=0)
print(x_train.shape)
print(x_test.shape)

(8000, 13)
```

(2000, 13)

Scale the independent variables

```
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.transform(x_test)
joblib.dump(sc,"churnsc.pkl")
['churnsc.pkl']
```